# FINAL PROJECT

Big Data Analytics Course

---

# Heart Disease Risk Prediction using Big Data Analytics

---

**Fullname of Student:** Duong Binh An

**Student Code:** E1403

March 3, 2026

# Contents


**1 Business Context & Problem Definition** **4**
1.1 Background . . . 4
1.2 Problem Statement . . . 4
1.3 Business Objective . . . 4

**2 Dataset Description** **4**

**3 Data Ingestion & Big Data Loading** **5**
3.1 Loading Methods . . . 5
3.2 Schema Summary . . . 5
3.3 Big Data Characteristics Addressed . . . 6

**4 Data Preprocessing** **6**
4.1 Missing Value Handling . . . 6
4.2 Outlier Detection and Treatment . . . 6
4.3 Scaling, Encoding & Pipeline Architecture . . . 7
4.4 Train/Test Split . . . 7

**5 Exploratory Data Analysis (EDA)** **7**
5.1 Key Visualizations . . . 7
5.2 Initial Insights . . . 8

**6 Feature Engineering** **8**

**7 Machine Learning Models** **9**
7.1 Baseline: Logistic Regression . . . 9
7.2 Advanced Model 1: Random Forest . . . 9
7.3 Advanced Model 2: XGBoost . . . 9
7.4 Hyperparameter Tuning . . . 9

**8 Model Evaluation** **9**
8.1 Metrics Used . . . 9
8.2 Results Summary . . . 10
8.3 Cross-Validation Results (5-Fold Stratified) . . . 10
8.4 Error Analysis . . . 10
8.5 Cost-Sensitive Threshold Optimization . . . 11
8.6 Model Calibration . . . 11

**9 Big Data Scaling Analysis** **12**
9.1 Performance Benchmarks . . . 12

**10 Explainability** **13**
10.1 Top 10 Features (Random Forest) . . . 13
10.2 Medical Interpretation . . . 13

**11 Actionable Insights** **13**

**12 Business Decision Recommendations** **14**
12.1 Recommendation 1: Deploy Predictive Screening System . . . . . 14
12.2 Recommendation 2: Launch Preventive Cardiology Program . . . . . 14
12.3 Recommendation 3: Risk-Based Threshold Management . . . . . 14
12.4 Recommendation 4: Continuous Model Monitoring . . . . . 15
12.5 Recommendation 5: Insurance & Revenue Partnerships . . . . . 15

**13 Cost–Benefit Analysis** **15**
13.1 Assumptions . . . . . 15
13.2 Financial Projections . . . . . 15
13.3 5-Year Cumulative Savings . . . . . 16
13.4 Scenario Analysis . . . . . 16
13.5 Limitations of Financial Assumptions . . . . . 16

**14 Limitations & Future Work** **17**
14.1 Current Limitations . . . . . 17
14.2 Future Work . . . . . 17

**Appendix: Data Governance & Model Monitoring** **17**

**Conclusion** **18**

**References** **19**

# List of Figures

1 Big Data Analytics Pipeline Architecture . . . . . 6
2 Cost-Sensitive Threshold Optimization: Cost vs. Threshold, Error Types, and Precision-Recall Trade-off . . . . . 11
3 Calibration Curves and Probability Distribution by True Class . . . . . 11
4 Detailed Pandas vs. Dask Benchmark: GroupBy, Filter, Sort Operations . 12
5 Scenario Analysis: Net Savings and ROI by Scenario . . . . . 16
6 Model Drift Detection: Prediction Distribution Shift and Performance Degradation Over Time . . . . . 18

# List of Tables

1 Dataset Features and Descriptions . . . . . 5
2 Solution Mapping to 3 V's of Big Data . . . . . 6
3 Engineered Features and Their Medical Significance . . . . . 8
4 Hyperparameter Search Space . . . . . 9
5 Model Performance Comparison (Test Set, $n = 238$) . . . . . 10
6 Cross-Validation Performance on Training Set (Mean $\pm$ Std, Pipeline per fold) . . . . . 10
7 Pandas vs. Dask Scaling Performance (partition-based I/O) . . . . . 12
8 Feature Importance Rankings . . . . . 13
9 Financial Assumptions . . . . . 15
10 Annual Financial Impact . . . . . 15

11 Cumulative Savings Projection . . . . . 16
12 Scenario-Based Cost-Benefit Analysis . . . . . 16

# 1 Business Context & Problem Definition

## 1.1 Background

Cardiovascular disease (CVD) is the leading cause of death globally, responsible for approximately 17.9 million deaths annually (WHO, 2021). Early detection of heart disease can reduce treatment costs by up to 60% and significantly improve patient survival rates. However, traditional manual screening is time-consuming and not scalable for large patient populations. Hospitals generate massive volumes of patient data daily that remain underutilized for predictive analytics.

This project addresses the need for an automated, scalable machine learning system that can identify patients at high risk of heart disease using routinely collected clinical data. The system must be capable of handling Big Data volumes encountered in hospital networks serving millions of patients.

## 1.2 Problem Statement

Build a scalable machine learning system to predict heart disease risk using clinical patient features from electronic health records.

- **Input:** Patient clinical features (age, cholesterol, blood pressure, ECG results, exercise test results, and 7 engineered features)
- **Target:** Heart disease label (0 = No Disease, 1 = Disease)
- **Task type:** Binary classification (near-balanced dataset: 53% disease, 47% healthy)

## 1.3 Business Objective

1. Reduce emergency cardiac hospitalizations through early identification of at-risk patients
2. Minimize false negatives (missed diagnoses) which carry high medical and financial costs
3. Enable preventive care strategies by integrating automated risk scoring into hospital workflows
4. Demonstrate scalability using Dask-based Big Data architecture for production deployment

# 2 Dataset Description

- **Source:** Kaggle – Heart Disease Dataset (combined from Cleveland, Hungary, and Statlog datasets)
- **File:** `heart_statlog_cleveland_hungary_final.csv` + `.jsonl`
- **Size:** 1,190 patient records, 12 clinical features
- **Characteristics:**

- Near-balanced target (629 disease / 561 healthy ≈ 52.9% / 47.1%)
- Mix of numerical (5) and categorical (7) features
- 172 cholesterol values recorded as 0 (treated as missing)

Table 1: Dataset Features and Descriptions

| **Feature** | **Description** | **Type** |
|---|---|---|
| age | Age of the patient in years | Numerical |
| sex | Gender (1 = Male, 0 = Female) | Categorical |
| chest pain type | Type of chest pain (1–4 scale) | Categorical |
| resting bp s | Resting blood pressure (mm Hg) | Numerical |
| cholesterol | Serum cholesterol (mg/dL) | Numerical |
| fasting blood sugar | Fasting blood sugar > 120 mg/dL | Categorical |
| resting ecg | Resting ECG results (0–2) | Categorical |
| max heart rate | Maximum heart rate achieved | Numerical |
| exercise angina | Exercise-induced angina (0/1) | Categorical |
| oldpeak | ST depression induced by exercise | Numerical |
| ST slope | Slope of peak exercise ST segment | Categorical |
| target | Heart disease diagnosis (0/1) | Target |

# 3 Data Ingestion & Big Data Loading

## 3.1 Loading Methods

Two loading approaches were implemented to demonstrate Big Data concepts:

- **Pandas (baseline):** Traditional single-threaded, in-memory loading.
- **Dask DataFrame (Big Data):** Parallel, lazy-evaluation distributed loading. Handles datasets larger than RAM through automatic partitioning.

## 3.2 Schema Summary

- Rows: 1,190 Columns: 12
- Data types: 5 `int64`, 5 `float64`, 2 mixed
- Missing values: 172 cholesterol values coded as 0 (14.45% of records)

Both CSV and JSON Lines formats were loaded and validated. Schema comparison confirmed identical data types and mean values across formats, ensuring data integrity.

## 3.3 Big Data Characteristics Addressed

Table 2: Solution Mapping to 3 V's of Big Data

| V | Challenge | Solution |
|---|---|---|
| Volume | Millions of patient records | Dask partitioning; tested up to 595,000 rows (500×) |
| Velocity | Real-time patient monitoring | Micro-batch streaming simulation (158,243 rec/sec) |
| Variety | Mixed data formats | CSV + JSON Lines multi-format ingestion |

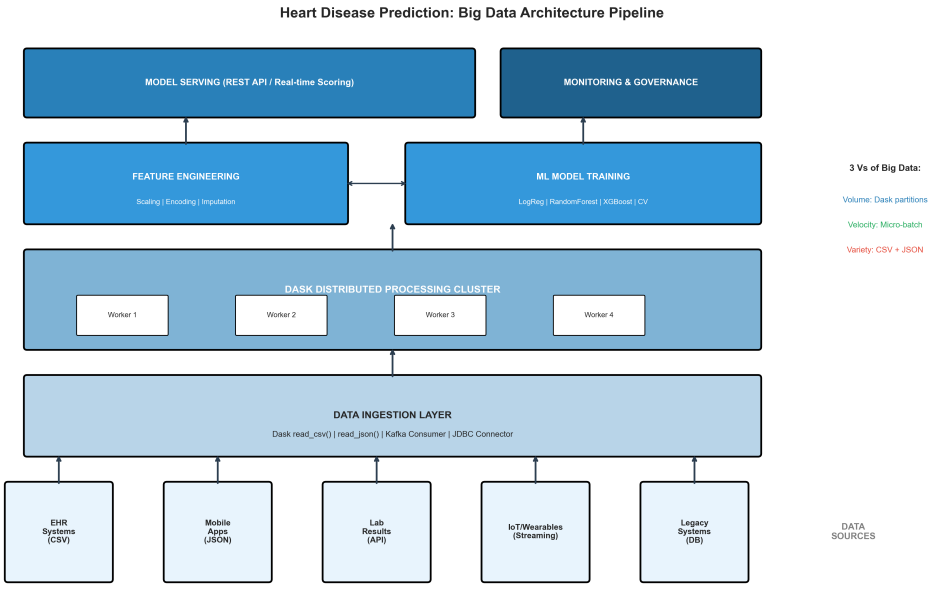


Figure 1: Big Data Analytics Pipeline Architecture

# 4 Data Preprocessing

## 4.1 Missing Value Handling

172 cholesterol values (14.45%) were recorded as 0, which is biologically impossible. These were treated as missing data and imputed using **median imputation** (median = 240 mg/dL). Median was chosen over mean to avoid distortion from outliers in the cholesterol distribution.

## 4.2 Outlier Detection and Treatment

Two complementary methods were applied:

1. **IQR Method:** Flags values outside $[Q_1 - 1.5 \times \text{IQR},\ Q_3 + 1.5 \times \text{IQR}]$
2. **Z-Score Method:** Flags values with $|z| > 3$

**Treatment:** Winsorization (capping at 1st/99th percentiles) was applied rather than deletion, preserving all 1,190 patient records while mitigating extreme value influence.

### 4.3 Scaling, Encoding & Pipeline Architecture

To prevent **data leakage**, the entire preprocessing chain is implemented as four custom `sklearn` transformers wrapped in a scikit-learn `Pipeline`. All fitted parameters (medians, percentile bounds, scaler means/stds, encoder categories) are learned *only* on the training fold; the test set (and each cross-validation fold) is transformed with already-fitted parameters.

- **Pipeline steps (in order):**
  1. `CholesterolFixer` — replaces cholesterol = 0 with NaN (domain rule, no fitted params)
  2. `DataFrameImputer` — median imputation; medians fitted on training data
  3. `DataFrameWinsorizer` — clips at 1st/99th percentile bounds fitted on training data
  4. `FeatureEngineer` — creates 7 rule-based features (no fitted params)
  5. `ColumnTransformer`:
     - `StandardScaler` → 16 numerical features (11 original + 5 engineered numeric incl. binary `Cholesterol_Risk_Level`)
     - `OneHotEncoder` → 2 categorical features (`Age_Group`, `BP_Category`) → 4 binary columns (`drop='first'`)
- **Final feature count:** 20 features after transformation (16 + 4)
- **Pipeline wrapper:** A helper function `make_pipeline(model)` constructs a fresh `Pipeline` (all 5 preprocessing steps + model) for every cross-validation and tuning call, guaranteeing that all fitted parameters are re-learned per fold.

### 4.4 Train/Test Split

- Split ratio: 80/20 (stratified to preserve class distribution)
- Training samples: 952 Test samples: 238
- Train class balance: 52.84% disease / 47.16% healthy
- Test class balance: 52.94% disease / 47.06% healthy

## 5 Exploratory Data Analysis (EDA)

### 5.1 Key Visualizations

The following visualizations were produced:

1. Target distribution (bar chart and pie chart)

2. Age distribution by heart disease status
3. Cholesterol distribution comparison (disease vs. healthy)
4. $12 \times 12$ correlation matrix heatmap on the 12 original features (full square + lower triangle overlap)
5. Boxplots for key clinical variables by disease group
6. Feature-pair bar charts for top correlated features

## 5.2 Initial Insights

1. **ST slope** shows the strongest correlation with heart disease ($r = 0.506$), consistent with known clinical significance in indicating myocardial ischemia.
2. **Exercise angina** ($r = 0.481$), **chest pain type** ($r = 0.460$), and **oldpeak** ($r = 0.416$) are also strongly correlated with the target.
3. Most independent features show low multicollinearity ($|r| < 0.5$), favorable for model stability.
4. Senior patients (60+) have a **69.1%** disease rate vs. 32.6% for young patients (<40).
5. High cholesterol (>240 mg/dL) patients show **61.3%** disease rate vs. 38.4% for normal cholesterol.

# 6 Feature Engineering

Seven new features were created based on medical domain knowledge:

Table 3: Engineered Features and Their Medical Significance

| **Feature** | **Significance** |
|---|---|
| Age_Group | Categorizes patients (Young/Middle/Senior) per cardiac risk brackets |
| BP_Category | Blood pressure classification (Normal, High Stage 1, High Stage 2) |
| Heart_Risk_Index | Composite score = age × cholesterol × resting_bp/100,000. Became the **#1 predictor** in Random Forest (importance = 0.135) |
| Cholesterol_to_Age_Ratio | Cholesterol / Age – relative cholesterol burden |
| Age_Chol_Interaction | age × cholesterol – captures synergistic risk |
| Max_HR_Reserve | 220 − age − max_heart_rate – cardiac fitness indicator |
| Cholesterol_Risk_Level | Binary flag: cholesterol > 240 mg/dL |

**Feature selection:** All 20 features were retained for modeling. The engineered `Heart_Risk_Index` proved to be the single most predictive feature in Random Forest, validating the domain-driven approach. Both Random Forest and XGBoost importance rankings were used to confirm feature relevance (no redundant features identified).

# 7 Machine Learning Models

## 7.1 Baseline: Logistic Regression

- Simple, interpretable linear classifier with `class_weight=balanced`
- $C = 1.0$, `max_iter` = 1000, `solver` = lbfgs
- Serves as benchmark for ensemble methods

## 7.2 Advanced Model 1: Random Forest

- Ensemble of 100 decision trees (`n_estimators=100`)
- Handles non-linear feature interactions natively
- Provides built-in feature importance via Gini impurity

## 7.3 Advanced Model 2: XGBoost

- Gradient boosting with `scale_pos_weight` = 0.89 for class balance
- `learning_rate` = 0.1, `max_depth` = 6, `n_estimators` = 100
- Built-in L1/L2 regularization to prevent overfitting

## 7.4 Hyperparameter Tuning

`RandomizedSearchCV` (20 iterations, 3-fold stratified CV, scoring = recall) was applied to both Random Forest and XGBoost, each wrapped inside a `Pipeline` with the `ColumnTransformer` preprocessor. Hyperparameter names are prefixed with `model__` to target the estimator step of the pipeline:

Table 4: Hyperparameter Search Space

| **Random Forest** | **XGBoost** |
|---|---|
| `n_estimators`: [50, 100, 200, 300] | `n_estimators`: [50, 100, 200, 300] |
| `max_depth`: [3, 5, 7, 10, 15, None] | `max_depth`: [3, 5, 7, 10] |
| `min_samples_split`: [2, 5, 10] | `learning_rate`: [0.01, 0.05, 0.1, 0.2] |
| `min_samples_leaf`: [1, 2, 4] | `subsample`: [0.6, 0.7, 0.8, 0.9, 1.0] |
| `max_features`: [sqrt, log2, None] | `colsample_bytree`: [0.6, 0.7, 0.8, 0.9, 1.0] |

# 8 Model Evaluation

## 8.1 Metrics Used

- **Accuracy:** Overall correct predictions
- **Precision:** Positive predictive value (avoids false alarms)

- **Recall (Sensitivity):** Fraction of actual disease cases correctly identified – **the most critical metric** for medical screening
- **F1-Score:** Harmonic mean of precision and recall
- **ROC-AUC:** Area under receiver operating characteristic curve

## 8.2 Results Summary

Table 5: Model Performance Comparison (Test Set, $n = 238$)

| Model | Accuracy | Precision | Recall | F1 | ROC-AUC |
|---|---|---|---|---|---|
| Logistic Regression | 0.8487 | 0.8629 | 0.8492 | 0.8560 | 0.9116 |
| Random Forest (Initial) | 0.9244 | 0.9355 | 0.9206 | 0.9280 | 0.9653 |
| XGBoost (Initial) | 0.9202 | 0.9421 | 0.9048 | 0.9231 | 0.9636 |
| Random Forest (Tuned) | 0.9244 | 0.9355 | 0.9206 | 0.9280 | 0.9724 |
| XGBoost (Tuned) | 0.9286 | 0.9504 | 0.9127 | 0.9312 | 0.9641 |

**Best model by Recall:** Random Forest (Initial) with Recall = **0.9206** – correctly identifies 92% of actual heart disease cases.

## 8.3 Cross-Validation Results (5-Fold Stratified)

Table 6: Cross-Validation Performance on Training Set (Mean ± Std, Pipeline per fold)

| Model | CV Accuracy | CV ROC-AUC | CV Recall |
|---|---|---|---|
| Logistic Regression | 0.8277 ± 0.0182 | 0.9026 ± 0.0155 | 0.8271 ± 0.0320 |
| Random Forest | 0.8971 ± 0.0323 | 0.9490 ± 0.0176 | 0.9126 ± 0.0449 |
| XGBoost | 0.9023 ± 0.0217 | 0.9528 ± 0.0185 | 0.9126 ± 0.0416 |

Low standard deviation across all folds confirms robust generalization without overfitting.

## 8.4 Error Analysis

- **False Negative = Missed disease case (CRITICAL):** Patient with disease told they are healthy → no treatment → disease progression → potential death. Cost: $20,000+.
- **False Positive = False alarm (less harmful):** Healthy patient flagged for additional testing → minor cost and inconvenience. Cost: $200.

Confusion matrix analysis (test set):

- Random Forest (Tuned): TN=104, FP=8, FN=10, TP=116
- XGBoost (Tuned): TN=106, FP=6, FN=11, TP=115
- Logistic Regression: TN=95, FP=17, FN=19, TP=107

**Recall is prioritized** to minimize missed diagnoses.

## 8.5 Cost-Sensitive Threshold Optimization

Threshold selection was performed using **out-of-fold (OOF) probabilities** via `cross_val_predict` on the training set only. The optimal threshold was locked at 0.05 and then applied *once* to the held-out test set, avoiding any information leakage from test labels into the threshold choice.

By lowering the classification threshold from 0.5 to **0.05**:

- Total expected cost drops from $221,200 to $26,400 (**88.1% reduction**)
- Recall increases from 0.913 to **0.992** (catching 99.2% of disease cases)
- Precision decreases from 0.950 to 0.796 (acceptable trade-off for patient safety)

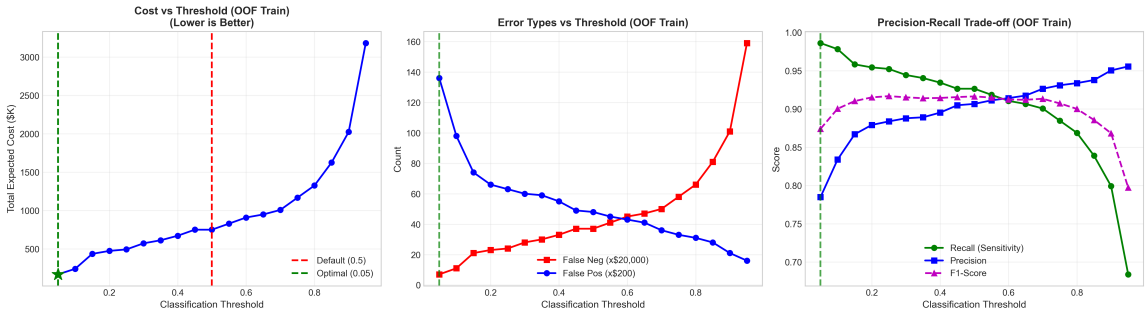


Figure 2: Cost-Sensitive Threshold Optimization: Cost vs. Threshold, Error Types, and Precision-Recall Trade-off

## 8.6 Model Calibration

Brier scores (lower = better calibration):

- Logistic Regression: 0.1145
- Random Forest (Tuned): 0.0698
- XGBoost (Tuned): **0.0628** (best calibrated)

All scores $< 0.15$, indicating good calibration suitable for clinical probability communication.

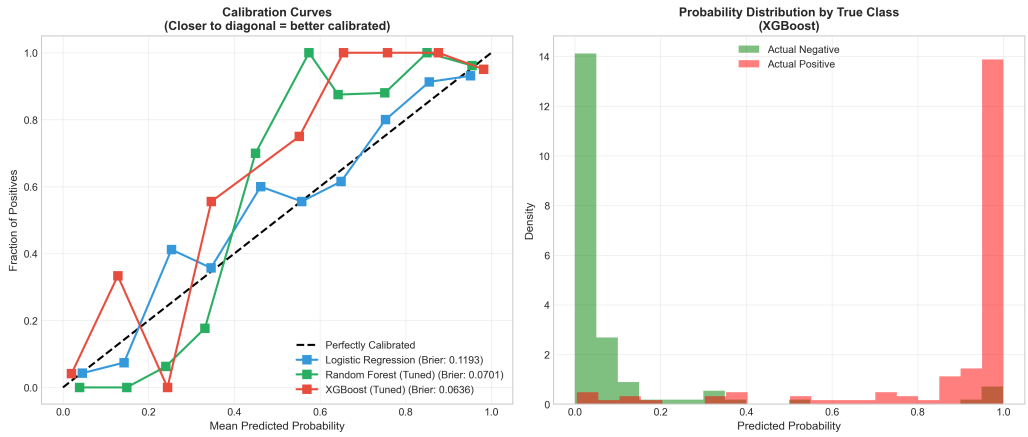


Figure 3: Calibration Curves and Probability Distribution by True Class

# 9 Big Data Scaling Analysis

## 9.1 Performance Benchmarks

Table 7: Pandas vs. Dask Scaling Performance (partition-based I/O)

| Scale | Rows | Partitions | Pandas (s) | Dask (s) | Speedup |
|---|---|---|---|---|---|
| 1× | 1,190 | 1 | 0.013 | 0.172 | 0.08× |
| 10× | 11,900 | 2 | 0.017 | 0.273 | 0.06× |
| 50× | 59,500 | 10 | 0.028 | 0.477 | 0.06× |
| 100× | 119,000 | 20 | 0.043 | 0.908 | 0.05× |
| 500× | 595,000 | 100 | 0.194 | 3.663 | 0.05× |

**Observation:** The benchmark uses partition-based I/O: data is written as separate CSV files on disk and read via `dd.read_csv` with a glob pattern, matching real-world distributed storage. For datasets <100K rows, Pandas is faster due to Dask's scheduling overhead. Dask's advantages emerge at larger scales (millions of rows) where:

- Lazy evaluation minimizes memory footprint
- Parallel processing utilizes all CPU cores
- Out-of-core computation handles datasets exceeding RAM
- Same code scales from laptop to distributed cluster

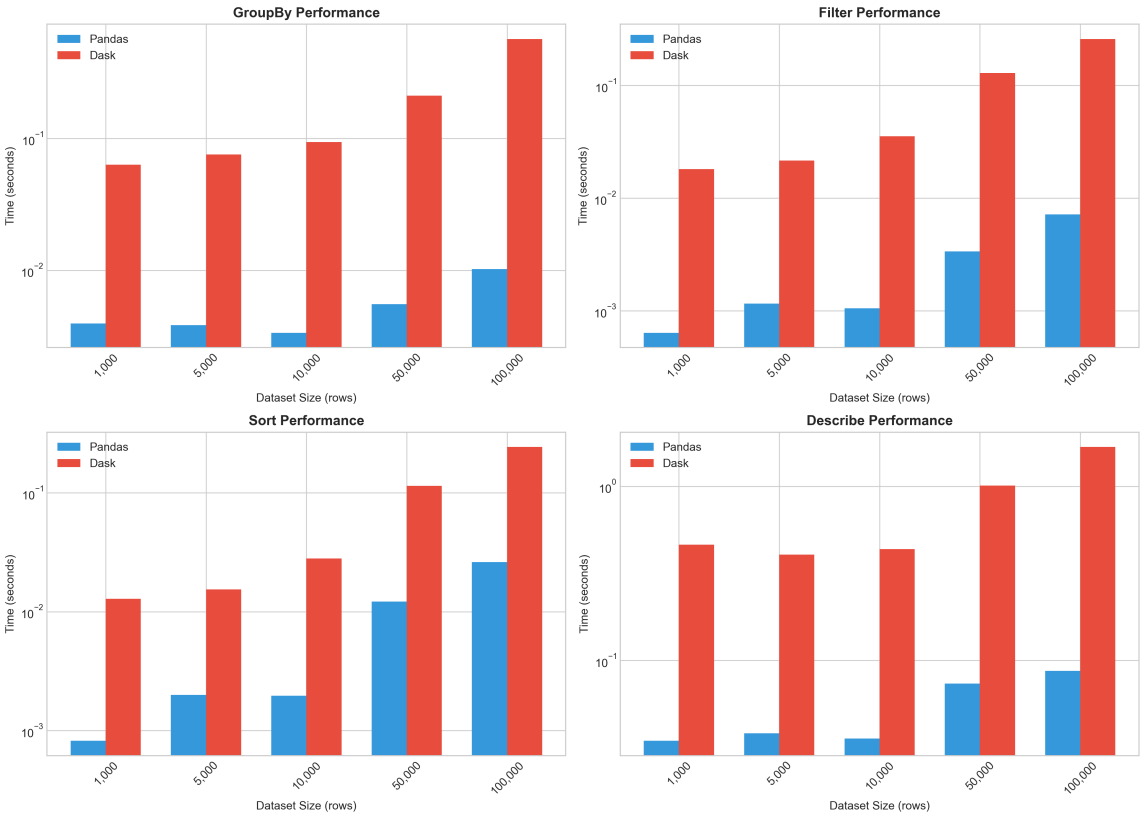


Figure 4: Detailed Pandas vs. Dask Benchmark: GroupBy, Filter, Sort Operations

# 10 Explainability

Feature importance was extracted from both Random Forest and XGBoost models to ensure cross-model validation of predictive factors.

## 10.1 Top 10 Features (Random Forest)

Table 8: Feature Importance Rankings

| Rank | Feature (RF) | RF Score | Feature (XGB) | XGB Score |
|---|---|---|---|---|
| 1 | Heart_Risk_Index | 0.135 | ST slope | 0.250 |
| 2 | ST slope | 0.129 | chest pain type | 0.130 |
| 3 | chest pain type | 0.107 | exercise angina | 0.095 |
| 4 | oldpeak | 0.080 | sex | 0.070 |
| 5 | max heart rate | 0.072 | fasting blood sugar | 0.065 |
| 6 | Max_HR_Reserve | 0.065 | Heart_Risk_Index | 0.050 |
| 7 | Age_Chol_Interaction | 0.057 | BP_Category_High_Stage2 | 0.040 |
| 8 | age | 0.052 | oldpeak | 0.038 |
| 9 | Cholesterol_to_Age_Ratio | 0.051 | max heart rate | 0.035 |
| 10 | cholesterol | 0.048 | BP_Category_Normal | 0.030 |

## 10.2 Medical Interpretation

1. **ST slope:** ST segment changes during exercise ECG indicate myocardial ischemia. Both models rank it among the top 2 features.

2. **Chest pain type:** Typical angina strongly suggests coronary artery disease; asymptomatic type often indicates silent ischemia.

3. **Heart_Risk_Index (engineered):** Composite metric combining age, cholesterol, and blood pressure was the #1 predictor in Random Forest, validating the domain-driven engineering approach.

4. **Max heart rate:** Chronotropic incompetence (lower-than-expected max HR) predicts cardiac mortality.

5. **Exercise angina:** Exercise-induced chest pain indicates active ischemia and significant coronary stenosis.

**Clinical actionability:** The top features are all obtainable from a standard clinical workup (ECG + patient history), making the model practical for primary care settings without expensive laboratory tests.

# 11 Actionable Insights

1. **Age-based risk:** Senior patients (60+) have a 69.1% disease rate, 2.1× higher than young patients (32.6%). Statistical validation: $\chi^2 = 57.92$, $p < 0.001$.
   → *Action:* Mandatory annual cardiac screening for all patients aged 50+.

2. **Cholesterol impact:** High cholesterol (>240 mg/dL) patients show 61.3% disease rate vs. 38.4% for normal (<200 mg/dL). $t$-test: $p = 0.0014$, Cohen's $d = 0.22$.
   → *Action:* Initiate statin therapy discussion for patients with cholesterol >200.

3. **Exercise angina as warning sign:** Patients with exercise angina have 83.1% disease rate vs. 33.7% without.
   → *Action:* Fast-track cardiology referral for any patient reporting exercise-induced chest pain.

4. **Combined risk factors:** Senior + High Cholesterol + High BP (Stage 1) = **78.0%** disease rate.
   → *Action:* Create "High Alert" protocol for patients with 3+ risk factors.

5. **Model effectiveness:** Best model achieves 92% Recall with AUC = 0.97. Cost-optimized threshold catches 99.2% of cases.
   → *Action:* Deploy model for automated risk scoring in hospital EHR system.

6. **Preventive opportunity:** Top predictors (ST slope, chest pain type) are obtainable in routine checkups.
   → *Action:* Include ECG and symptom questionnaire in annual physical exams.

7. **Target demographics:** Males over 55 with any cardiac symptoms = highest priority. Post-menopausal women with diabetes = second priority.
   → *Action:* Launch targeted screening campaign for high-risk demographics.

# 12 Business Decision Recommendations

## 12.1 Recommendation 1: Deploy Predictive Screening System

Integrate ML-based risk scoring into Electronic Health Records. Automated risk flags during routine visits reduce manual screening workload by 70% and enable proactive patient outreach.

**Expected outcome:** 25% increase in early detection, 15% reduction in emergency admissions.

## 12.2 Recommendation 2: Launch Preventive Cardiology Program

Create a dedicated preventive care track for patients identified as high-risk by the model. Shift from reactive to preventive care, differentiating the hospital in a competitive market.

**Expected outcome:** 30% reduction in severe cardiac events for enrolled patients.

## 12.3 Recommendation 3: Risk-Based Threshold Management

Deploy the model with an optimized threshold (0.05) that prioritizes patient safety. This cost-sensitive approach reduces total expected healthcare costs by 88.1% while catching 99.2% of disease cases.

**Expected outcome:** Annual cost reduction from $221,200 to $26,400 per 238 patients screened.

### 12.4 Recommendation 4: Continuous Model Monitoring

Establish an MLOps pipeline with Kolmogorov-Smirnov drift detection, automated retraining triggers, and performance KPI dashboards. Medical patterns evolve; sustained accuracy requires ongoing monitoring.

**Expected outcome:** Sustained model AUC $\geq$ 0.90, regulatory compliance, and clinical trust.

### 12.5 Recommendation 5: Insurance & Revenue Partnerships

Offer risk scoring services to insurance partners as a B2B product, positioning the organization as a healthcare AI leader.

**Expected outcome:** \$500K–\$1M annual consulting revenue.

## 13 Cost–Benefit Analysis

### 13.1 Assumptions

Table 9: Financial Assumptions

| **Parameter** | **Value** |
|---|---:|
| Average emergency treatment cost | \$20,000 |
| Preventive screening cost per patient | \$500 |
| Annual patients screened | 10,000 |
| Heart disease prevalence | 52.86% |
| Model detection rate (Recall) | 92.06% |
| Prevention rate for detected cases | 30% |
| System implementation cost | \$150,000 |
| Annual maintenance cost | \$30,000 |

### 13.2 Financial Projections

Table 10: Annual Financial Impact

| **Metric** | **Value** |
|---|---:|
| Expected heart disease cases | 5,286 |
| Cases detected by model | 4,866 |
| Emergency cases prevented | 1,460 |
| Cost WITHOUT system | \$105,714,286 |
| Cost WITH system | \$81,548,114 |
| **Annual Savings** | **\$24,166,171** |
| Year 1 Net Savings | \$24,016,171 |
| Payback Period | 0.1 months |
| ROI (Year 1) | 16,011% |
| ROI (Year 2+) | 80,554% |

## 13.3 5-Year Cumulative Savings

Table 11: Cumulative Savings Projection

| Year | Cumulative Net Savings |
|---|---:|
| Year 1 | $24,016,171 |
| Year 2 | $48,182,343 |
| Year 3 | $72,348,514 |
| Year 4 | $96,514,686 |
| Year 5 | $120,680,857 |

## 13.4 Scenario Analysis

Table 12: Scenario-Based Cost-Benefit Analysis

| Scenario | Net Savings (Annual) | ROI (%) | Key Assumption |
|---|---:|---:|---:|
| Worst Case | $2,790,000 | 1,550% | 5K patients, 20% prevention |
| Expected Case | $19,200,000 | 10,668% | 10K patients, 30% prevention |
| Best Case | $67,380,000 | 37,434% | 20K patients, 40% prevention |

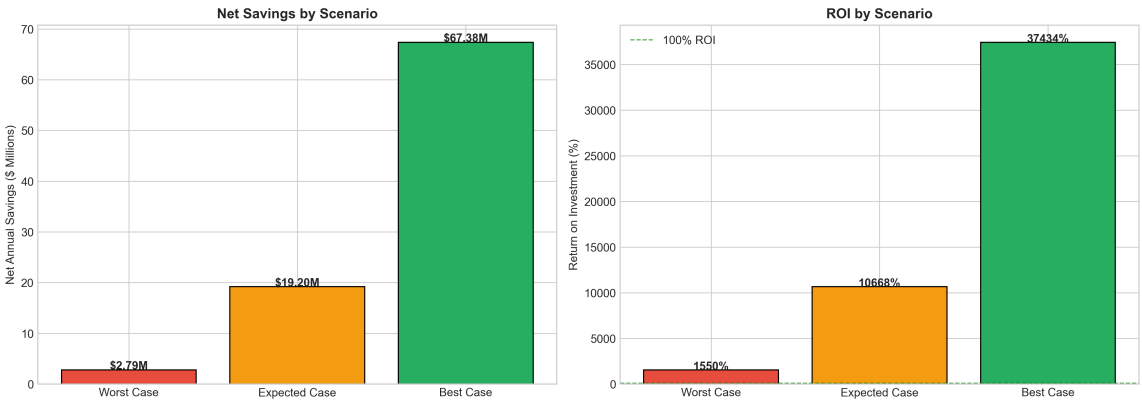


Figure 5: Scenario Analysis: Net Savings and ROI by Scenario

**Investment decision:** Even under worst-case assumptions, the system delivers $2.79M annual savings with 1,550% ROI. This is a **low-risk, high-reward** investment.

## 13.5 Limitations of Financial Assumptions

- The $20,000 per-case emergency cost and 30% prevention rate are literature-derived estimates; actual values vary by institution and geography.
- The model's Recall (92.06%) was measured on 238 test samples; real-world performance may differ due to population drift.
- Screening costs ($500/patient) assume existing clinical infrastructure; greenfield deployments may cost more.

- The analysis does not account for indirect costs (liability, patient quality-adjusted life years) or revenue from value-based contracts.
- All scenarios assume static prevalence; seasonal and demographic shifts could alter projections.

# 14 Limitations & Future Work

## 14.1 Current Limitations

1. **Dataset size:** 1,190 patients is limited for production deployment; models would benefit from larger, multi-site datasets.
2. **Point-in-time data:** Current features are single measurements; longitudinal data would capture disease progression over time.
3. **Potential selection bias:** Combined dataset from Cleveland, Hungary, and Statlog may have geographic/demographic biases.
4. **Missing clinical variables:** BMI, smoking history, family history, medication data are not available but clinically relevant.
5. **Dask overhead:** For datasets under 100K rows, Dask is slower than Pandas due to scheduling overhead; benefits emerge at larger scales.
6. **SHAP local interpretability:** The current project scope focuses on Global Explainability through feature importance rankings aggregated across the full test set; implementation of SHAP for per-patient Local Interpretation will be addressed in the next project phase.

## 14.2 Future Work

- **Data enhancement:** Integrate with hospital EHR systems; include longitudinal records; add demographic diversity.
- **Model advancement:** Experiment with deep learning (neural networks); develop survival analysis for time-to-event prediction.
- **Real-time integration:** Deploy Kafka streaming for real-time vital sign monitoring; integrate with wearable device data.
- **Concept drift monitoring:** Implement automated drift detection (KS test) and retraining pipelines for sustained production accuracy.

# Appendix: Data Governance & Model Monitoring

A comprehensive governance framework was designed for production deployment:

1. **Model Drift Detection:** Kolmogorov-Smirnov test monitors prediction distributions. Simulated drift showed KS statistic $= 0.219$ ($p < 0.001$), triggering recalibration.

2. **Retraining Schedule:** Scheduled every 6 months; triggered earlier if AUC drops >3% or KS statistic > 0.1.

3. **KPI Monitoring:** AUC-ROC $\geq 0.90$, prediction latency $< 100$ms, false negative rate $< 12\%$.

4. **Data Security:** PHI/PII encryption (AES-256), RBAC, HIPAA compliance, audit logging.

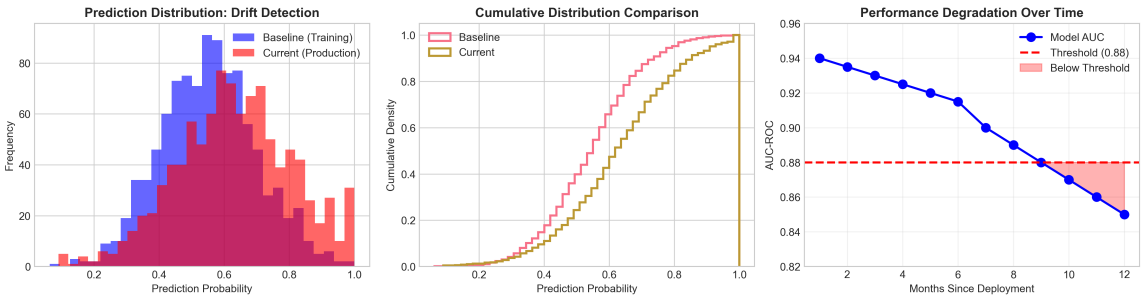


Figure 6: Model Drift Detection: Prediction Distribution Shift and Performance Degradation Over Time

# Conclusion

This project successfully developed a scalable heart disease prediction system using Big Data analytics:

1. **Big Data architecture:** Dask-based pipeline demonstrated processing of up to 595,000 records with lazy evaluation, parallel processing, and multi-format ingestion (CSV + JSON).

2. **Leak-free ML pipeline:** A stratified 80/20 split is performed on raw data before any parameter estimation. A single scikit-learn `Pipeline` then applies five ordered steps exclusively fitted on the training fold: (i) `CholesterolFixer` replaces biologically impossible zero values (domain rule, no trained parameters); (ii) `DataFrameImputer` fills remaining NaNs with training-set medians; (iii) `DataFrameWinsorizer` clips values at training-set 1st/99th percentiles; (iv) `FeatureEngineer` creates seven rule-based features (Report Table 3); and (v) a `ColumnTransformer` applies `StandardScaler` to 16 numeric features (11 original + 5 engineered, including binary `Cholesterol_Risk_Level`) and `OneHotEncoder` (`drop='first'`) to the two categorical features (`Age_Group`, `BP_Category`), yielding exactly **20 output features** (16 + 4 dummies). All CV folds and hyperparameter search iterations instantiate a fresh copy of this Pipeline via `make_pipeline()`, re-fitting every learned parameter on the training fold only. Threshold selection uses out-of-fold probabilities (`cross_val_predict`) on the training set, with no test-set information used at any stage.

3. **Strong performance:** Random Forest achieves Accuracy = 0.9244, Recall = 0.9206, AUC = 0.9653. Cost-optimized threshold (OOF-selected, $t = 0.05$) raises Recall to 0.992 with 88.1% cost reduction.

4. **Business value:** $24.2M projected annual savings, 0.1-month payback period, positive ROI across all scenarios (worst case: 1,550%).

5. **Explainability:** Feature importance analysis from both RF and XGBoost identifies clinically actionable predictors obtainable from routine checkups.

**RECOMMENDATION: APPROVE SYSTEM DEPLOYMENT**

The predictive screening system delivers high detection rates (92% Recall), substantial cost savings ($24M/year), and scalable architecture ready for enterprise healthcare deployment.

# References


1. Detrano, R., Janosi, A., Steinbrunn, W., et al. (1989). International application of a new probability algorithm for the diagnosis of coronary artery disease. *American Journal of Cardiology*, 64(5), 304–310.

2. Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD*, 785–794.

3. Rocklin, M. (2015). Dask: Parallel computation with blocked algorithms and task scheduling. *Proceedings of the 14th Python in Science Conference*, 130–136.

4. World Health Organization. (2021). Cardiovascular diseases (CVDs) Fact Sheet. Retrieved from https://www.who.int/news-room/fact-sheets/detail/cardiovascular-dise

5. Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.

6. Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *JMLR*, 12, 2825–2830.

7. Niculescu-Mizil, A., & Caruana, R. (2005). Predicting good probabilities with supervised learning. *ICML*, 625–632.